



**ClickHouse** in Action:  
Powering **Electrum**'s  
Electric Mobility  
Data Infrastructure

# Andi Pangeran



- Over **10 years of experience** in the IT Industry
- Currently, **Principal Software Engineer** at **Electrum.id**



@apangeran



@a\_pangeran



@A\_Pangeran



We aim to create a sustainable future where electric mobility is **accessible, efficient, and integral** to urban life. We offer electric motorcycles & mobility solutions with **convenient charging battery swap technology.**

**3,000+**

electric  
motorcycles

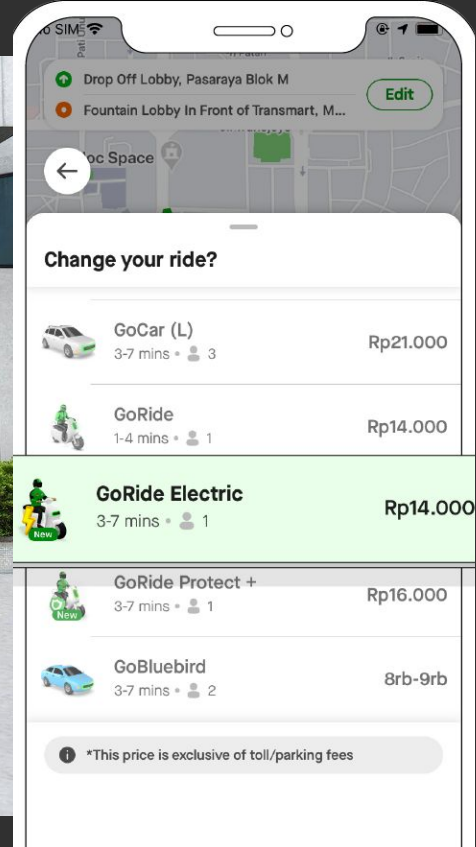
**250+**

swap stations  
across Jakarta

**220k**

km travelled  
per day





Drop Off Lobby, Pasaraya Blok M

Edit



### Change your ride?



GoCar (L)

3-7 mins • 3

Rp21.000



GoRide

1-4 mins • 1

Rp14.000



GoRide Electric

3-7 mins • 1

Rp14.000



GoRide Protect +

3-7 mins • 1

Rp16.000



GoBluebird

3-7 mins • 2

8rb-9rb



\*This price is exclusive of toll/parking fees



electrum 

### Paket Beli Baterai

Motor + Baterai + Charger

# Rp 23,9 Jt

OTR Jabodetabek



Battery



Portable Charger



DP Mulai 0%

Garansi Baterai  
5 tahun/  
50ribu KM

Cicilan Mulai 600 ribuan  
Tenor hingga 60 bulan

1+1 Tahun | 20.000 km  
Garansi Part\*

2+1 Tahun | 30.000 km  
Garansi Dynamo\*

1 Tahun | Gratis  
4x Service  
1x Part Service

1 Tahun | Gratis  
Home Services

Gratis  
Helm Electrum  
& Toolkit

1 Tahun | Emergency  
Road Assistance

electrum 

### Paket Sewa Baterai

Motor + Sewa Baterai + Swap

# Rp 19,9 Jt

OTR Jabodetabek



Swap Baterai  
di BSS Terdekat



~~Sewa Baterai Rp 230ribu/bulan~~  
**GRATIS Selama 18 Bulan**

DP Mulai 0%

Cicilan Mulai 500 ribuan  
Tenor hingga 60 bulan

1 Tahun | 10.000 km  
Garansi Part\*

2 Tahun | 30.000 km  
Garansi Dynamo\*

1 Tahun | Gratis  
Home Services

Additional  
Portable Charger  
Rp 1,2 Jt

Gratis  
Helm Electrum  
& Toolkit


\*Syarat dan ketentuan Berlaku

Pembelian Terbatas

\*Syarat dan ketentuan Berlaku


Pembelian Terbatas

titipan  
kantor



# Data at **Electrum**

# Data: *Internet of Things* (IoT)



**Battery**

**Battery Management System:**  
Voltage, current, temperature, state of charge (SOC), state of health (SOH), Cycle Count, Cells Information

**GPS:**  
GPS coordinates

**Vehicle:**  
Speed and distance traveled

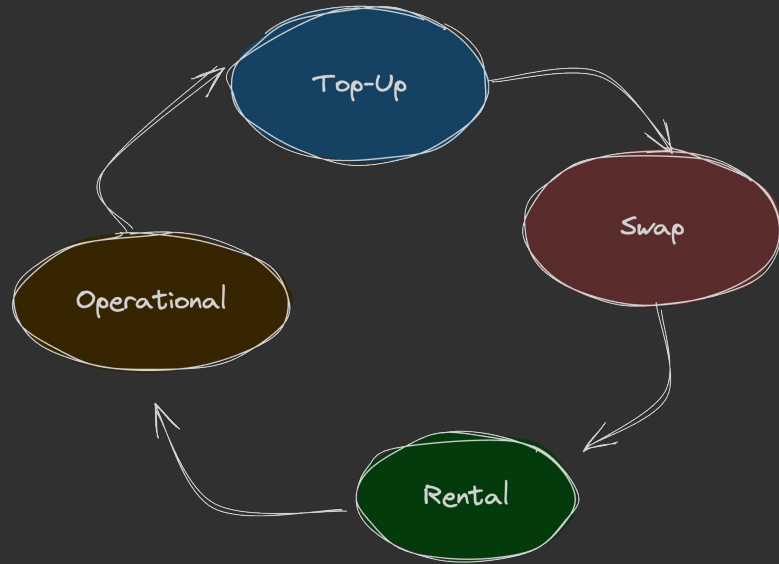


**Swap Station**

**General Data:**  
State, Input Power, voltage, current

**Doors Data:**  
State (lock/unlock), status (charging/idle), All batteries data

# Data: *Transactions* (Order)



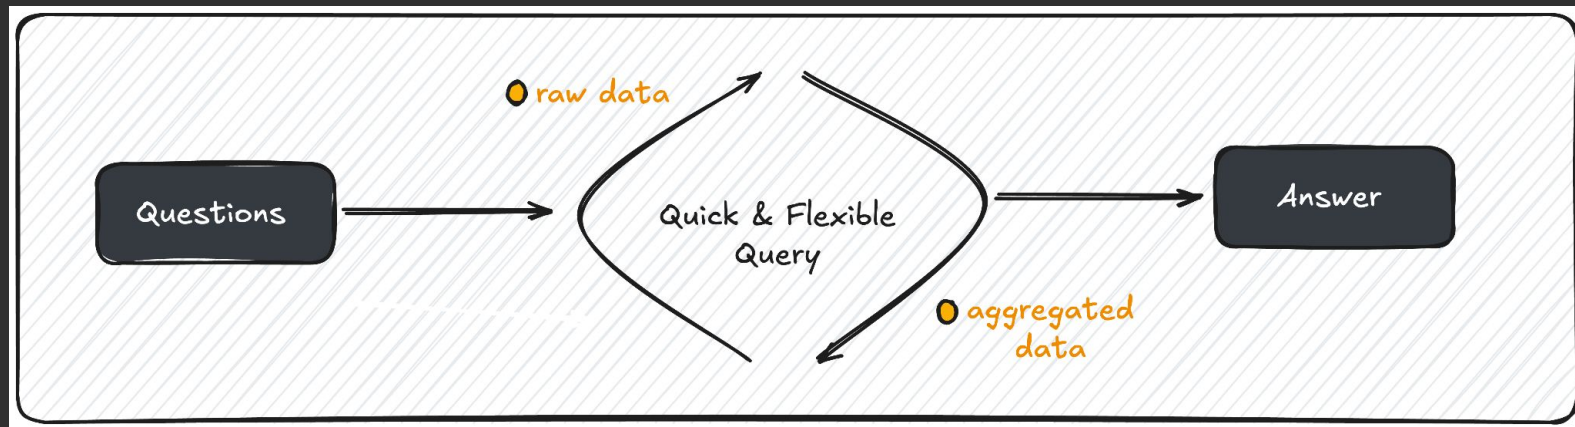
- **Top-Up Transactions:**  
Data on account recharges, including amounts, payment methods, and status.
- **Swap Transactions:**  
Details on battery swaps (station ID, door-number, battery ID, completion status).
- **Rental Transactions:**
  - Onboarding: Start data (user, vehicle, battery).
  - Offboarding: End data (user, vehicle, vehicle condition).
- **Operational Activity:**  
Maintenance, logistics, and fleet management data (e.g., technician actions, station repairs).



So what do we look for?

**Flexible, Scalable, & Low-Cost**  
Data Stack

# Need: *Quick & Flexible* Query



**Question:** Why do one of our battery swap stations was charging slower than others?

**Answer:** Query showed that the station's temperature is 4°C hotter than others, which caused slower charging to avoid overheating. Later, we learned the station was positioned under direct sunlight.

# Need: Scalable & Low-Cost

## Resource Efficiency

Efficiently use resources so we're not **over-provisioning**

## Expandable Infrastructure

Able to **scale linearly** with resources; as datasets grow, adding more resources will keep response times fast.



# Why ClickHouse?



Columnar  
Storage



Indexing &  
Data  
Skipping



Data  
Organization  
& Merging  
Mechanism

# Why: *It's columnar based*

completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price

**Row-based**

**Excess disk Reads** We need data in column, but it's stored in rows. So entire rows are read.

completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price
completed_at	order_number	order_state	order_type	total_price

**Column-Based**

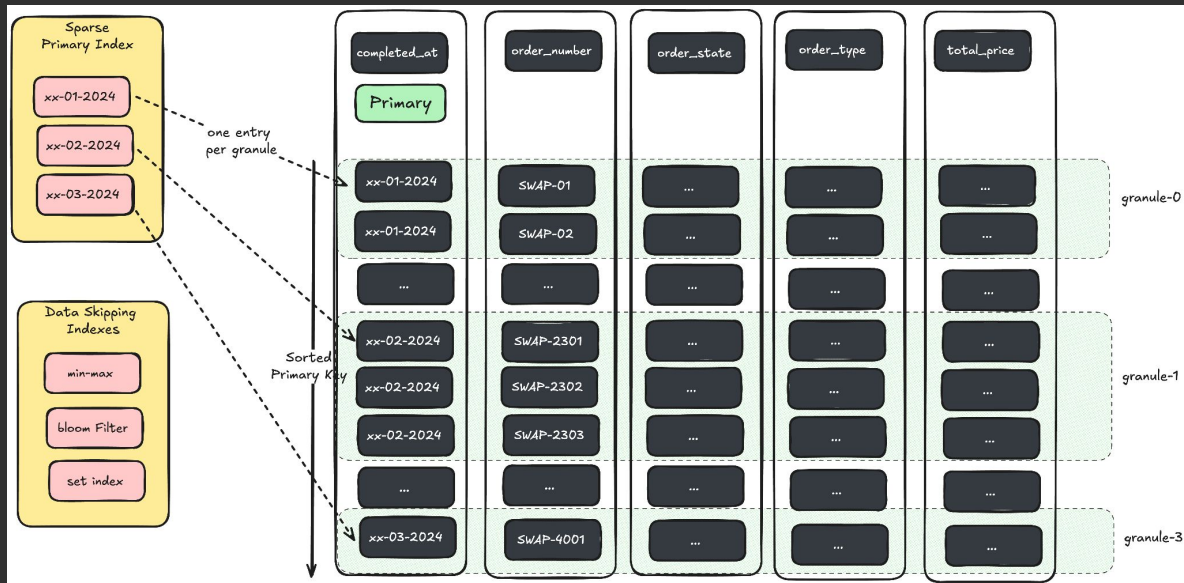
**Column Pruning** it only pick the subset of columns that are openly declared in the SQL query, thus eliminating 99% columns (990 of 1000 columns) from the disk read consideration.

## **Vectorized Execution**

data processed in **large arrays**, which could be the entire length of the column loaded into RAM.

+ **reduces CPU cache miss rates**

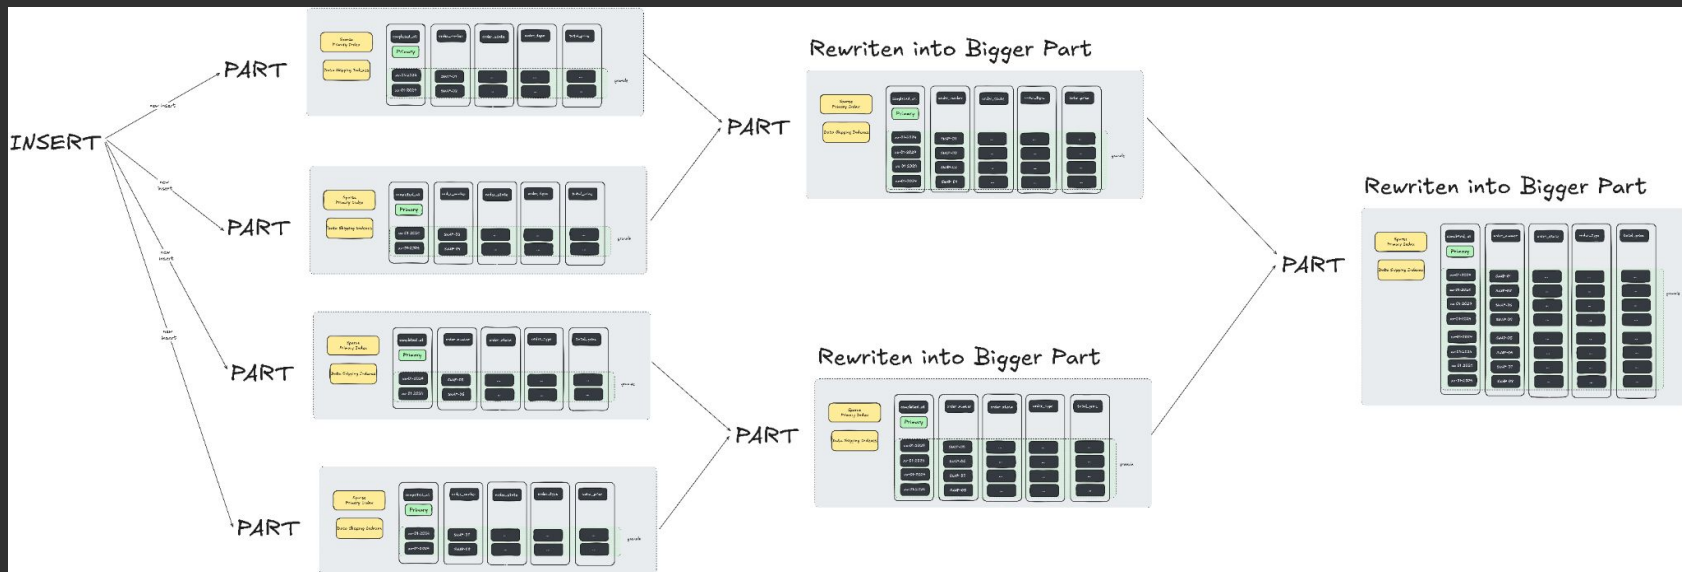
# Why: Multiple Layers of Indexes



- **Granules** are chunks of rows, and **ClickHouse** groups rows into granules based on the `index_granularity` setting.
- **Sparse Primary Index:** **ClickHouse** uses the `sparse_primary index` to quickly jump to the relevant granules, skipping over large portions of the data that don't need to be read.
- **Skip Indexes:** These allow **ClickHouse** to skip entire granules if it can determine from the index that no rows in the granule satisfy the query's filter.



# Why: *Efficient Merging*



**Parts:** Each time data is inserted, a new part is created. Multiple parts can accumulate over time, which can slow down query performance.

To improve performance and reduce the number of parts, **ClickHouse** periodically runs merging operations.

# What Happens During a Merge?




```
CREATE TABLE uk_price_paid
(
    date Date,
    town String,
    price UInt32
)
ENGINE = MergeTree
ORDER BY (town, date);
```

## MergeTree:

- Combines parts without deduplication or transformations.
- Simply reduces the number of parts but keeps all rows (including duplicates).
- Ideal for append-only data.

# What Happens During a Merge?



```
CREATE TABLE uk_listings
```

```
(  
  id UInt32,  
  date Date,  
  town String,  
  price UInt32  
)
```

```
ENGINE = ReplacingMergeTree(date)  
ORDER BY id;
```

## ReplacingMergeTree:

- Combines parts and deduplicates based on a key (or version).
- Replaces older rows with newer ones during a merge.
- Suitable for upserts and scenarios where data is frequently updated.

# What Happens During a Merge?



```
CREATE TABLE uk_price_paid_aggregates
(
  town String,
  max_price SimpleAggregateFunction(max, UInt32),
  avg_price AggregateFunction(avg, UInt32)
)
ENGINE = AggregatingMergeTree
ORDER BY town;
```

## AggregatingMergeTree:

- Combines parts by aggregating rows with the same key using predefined functions like sum, count, avg, etc.
- Ideal for pre-aggregating data to reduce query time on large datasets.

Another thing we love about ClickHouse is  
**materialized view.**

Let's deep-dive

# Materialized view 101

```
CREATE TABLE IF NOT EXISTS trx_orders
(
  completed_at DateTime64(3, 'UTC'),
  order_number String,
  order_type LowCardinality(String),
  total_price Int64
)

ENGINE = MergeTree
PARTITION BY toYYYYMM(completed_at)
ORDER BY (completed_at, order_type, order_number)
SETTINGS index_granularity = 8192;
```

**Raw-Table**

```
CREATE TABLE IF NOT EXISTS trx_agg_daily_orders
(
  order_date String,
  order_type LowCardinality(String),
  count_orders AggregateFunction(count),
  total_price AggregateFunction(sum, Int64)
)

ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(order_date)
ORDER BY (order_date, order_type);
```

**Aggregated-Table**

**SYNC ?**



# Materialized view 101

```
CREATE TABLE IF NOT EXISTS trx_orders
(
  completed_at DateTime64(3, 'UTC'),
  order_number String,
  order_type LowCardinality(String),
  total_price Int64
)
ENGINE = MergeTree
PARTITION BY toYYYYMM(completed_at)
ORDER BY (completed_at, order_type, order_number)
SETTINGS index_granularity = 8192;
```

```
CREATE TABLE IF NOT EXISTS trx_agg_daily_orders
(
  order_date String,
  order_type LowCardinality(String),
  count_orders AggregateFunction(count),
  total_price AggregateFunction(sum, Int64)
)
ENGINE = AggregatingMergeTree()
PARTITION BY toYYYYMM(order_date)
ORDER BY (order_date, order_type);
```

**Raw-Table** —————> **Materialized View** —————> **Aggregated-Table**

```
CREATE MATERIALIZED VIEW trx_agg_daily_orders_mv
TO trx_agg_daily_orders
AS

SELECT toDate(completed_at) AS order_date,
       order_type,
       countState()          AS total_orders,
       sumState(total_price) AS total_price
FROM trx_orders
GROUP BY order_date, order_type;
```

# Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256



**Raw-Table**



**Materialized View**



**Aggregated-Table**

# Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256



**Raw-Table**



**Materialized View**



**Aggregated-Table**

Insert new  
part on disk



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

# Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

**Raw-Table**

**Materialized View**

**Aggregated-Table**

Insert new  
part on disk

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

GROUP BY  
date, order-type

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

# Materialized view 101 - insert scenario empty

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

**Raw-Table**

Insert new  
part on disk

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

**Materialized View**

GROUP BY  
date, order-type

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

**Aggregated-Table**

STORE as new part  
of target table

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

# Materialized view 101 - insert merge scenario

**Raw-Table**     $\longrightarrow$     **Materialized View**     $\longrightarrow$     **Aggregated-Table**

Existing part  
on disk



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256

Existing part  
of target table



	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711



# Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624



**Raw-Table**



**Materialized View**



**Aggregated-Table**

Merge with  
existing part



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624



	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

# Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

**Raw-Table**

Merge with  
existing part

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

**Materialized View**

GROUP BY  
date, order-type

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	2	4016
2024-09-28	TOPUP	2	4958

**Aggregated-Table**

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

# Materialized view 101 - insert merge scenario

New insert block

completed_at	order_number	order_type	total_price
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

**Raw-Table**

Merge with  
existing part

completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

**Materialized View**

GROUP BY  
date, order-type

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	2	4016
2024-09-28	TOPUP	2	4958

**Aggregated-Table**

Merge with  
existing part  
Of target table

	order_date	order_type	total_orders	total_price
1	2024-09-28	SWAP	2	2339
2	2024-09-28	TOPUP	2	4711

completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	4	6355
2024-09-28	TOPUP	4	9669

# Materialized view 101 - Query Comparison

Raw-Table      —————>    Materialized View    —————>    Aggregated-Table



completed_at	order_number	order_type	total_price
2024-09-28 13:48:26.249280	ORD953128	SWAP	59
2024-09-28 13:48:26.249282	ORD957709	TOPUP	455
2024-09-28 13:48:26.249283	ORD475309	SWAP	2280
2024-09-28 13:48:26.249283	ORD110939	TOPUP	4256
2024-09-28 14:04:32.656515	ORD144868	SWAP	1392
2024-09-28 14:04:32.656517	ORD355138	TOPUP	623
2024-09-28 14:04:32.656517	ORD286434	TOPUP	4335
2024-09-28 14:04:32.656517	ORD847260	SWAP	2624

```
SELECT toDate(completed_at) AS order_date,  
       order_type,  
       count(*)              AS total_orders,  
       sum(total_price)      AS total_price  
FROM   trx_orders  
GROUP BY order_date, order_type;
```

**Processed 8 rows**



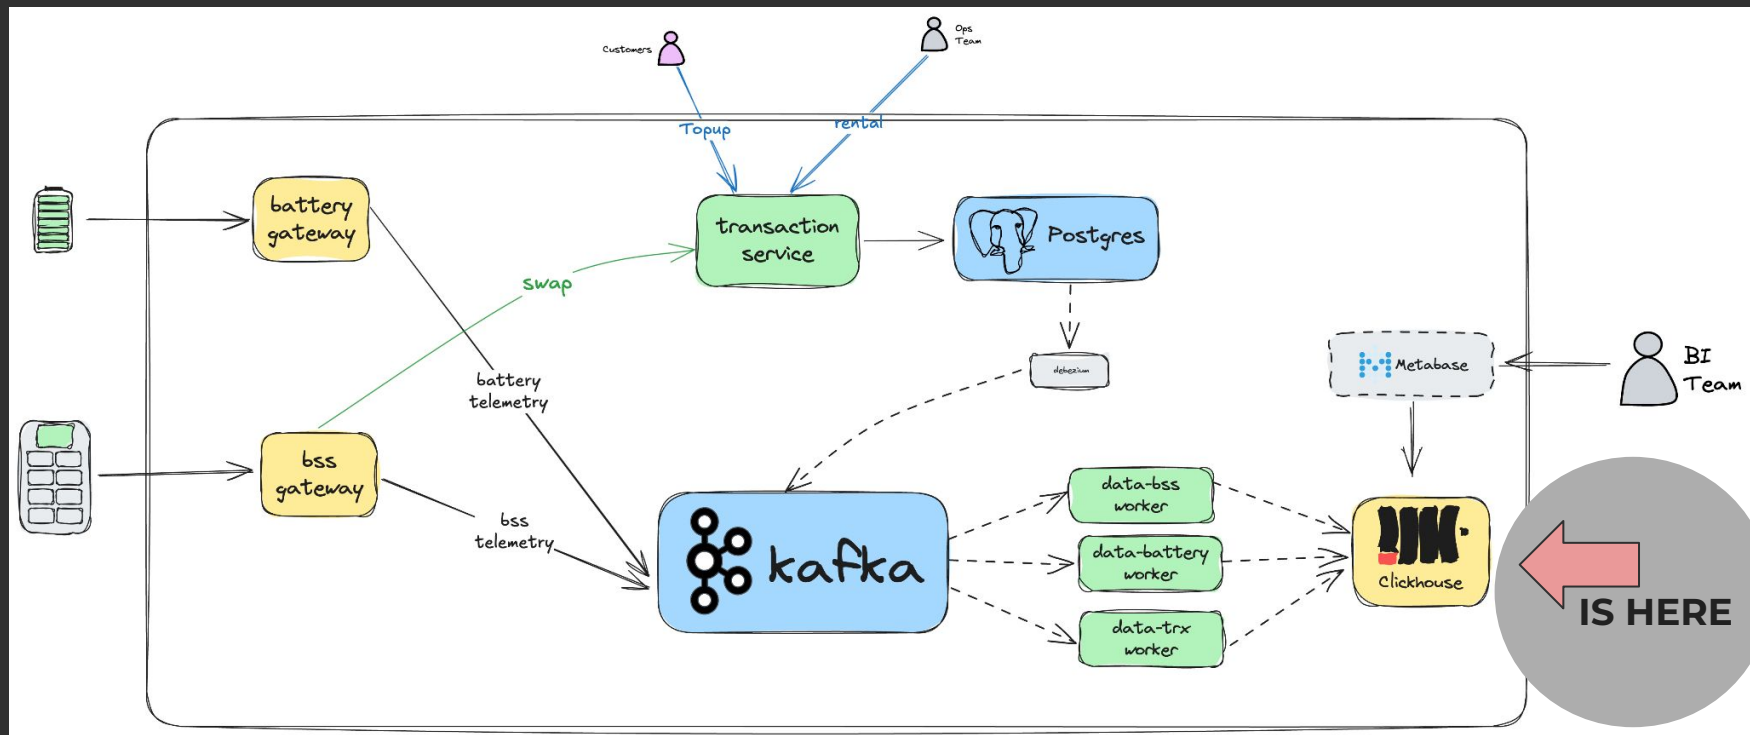
completed_at	order_type	total_orders	total_price
2024-09-28	SWAP	4	6355
2024-09-28	TOPUP	4	9669

```
SELECT  
  order date,  
  order type,  
  countState() AS total orders,  
  sumState(total price) AS total_price  
FROM   trx agg daily orders  
GROUP BY order_date, order_type;
```

**Processed 2 rows**

How ClickHouse  
**Powers** Our Data Platform

# How ClickHouse Powers Our Data Platform





# DEMO

<https://github.com/meong1234/clickhouse-meetup-10-2024>

# Where are we?

**50 million+**

Records / Day

**≤\$500/month**

*“**ClickHouse** enables us to have a data analytics capability that grows with our needs, while keeping costs low and predictable” - Electrum.id*

# Key Takeaways

## Columnar Format:

**ClickHouse** stores data in **columns**, optimizing read performance for analytical queries by reading only the relevant columns, reducing I/O and improving query speed.

## Sparse Primary Index:

**ClickHouse** uses a **sparse index** to store pointers to data blocks (granules), allowing it to **skip irrelevant data** during queries, improving performance when filtering by the primary key.

## Skip Index:

The **skip index** allows **ClickHouse** to **bypass granules** based on non-primary key column values, helping to avoid unnecessary data scanning and speeding up queries on non-primary key columns.

## Merging Process in Different Engines:

- **MergeTree**: Performs basic merging of data parts without deduplication or aggregation.
- **ReplacingMergeTree**: Deduplicates data during merging, keeping only the most recent version of rows with the same primary key.
- **AggregatingMergeTree**: Aggregates data during merging based on predefined aggregation functions, storing the results in summarized form.

## ClickHouse Materialized Views:

- **Triggered on Inserts**: Automatically updates the materialized view whenever new data is inserted into the source table.
- **Real-Time Pre-Aggregation**: Performs pre-aggregation or data transformation in real time, allowing for faster query performance by reducing the need for complex computations at query time.

# Ref

1. **Why ClickHouse Is So Fast**  
ChistaDATA Inc. (2023). Why ClickHouse Is So Fast.  
<https://chistadata.com/why-clickhouse-is-so-fast>
2. **How to Use ClickHouse Indexes: A Practical Guide**  
ChistaDATA Inc. (2023). How to Use ClickHouse Indexes: A Practical Guide.  
<https://chistadata.com/how-to-use-clickhouse-indexes-practical-guide>
3. **ClickHouse Skipping Indexes**  
ClickHouse Documentation. (2023). Skipping Indexes.  
<https://clickhouse.com/docs/en/optimize/skipping-indexes>
4. **ClickHouse Overview**  
YouTube. (2023). ClickHouse Overview.  
<https://www.youtube.com/watch?v=QDAJTKZT8y4>
5. **ClickHouse Architecture Deep Dive**  
YouTube. (2023). ClickHouse Architecture Deep Dive.  
<https://www.youtube.com/watch?v=iLXXoDaF0xs>
6. **Optimizing ClickHouse Performance**  
YouTube. (2023). Optimizing ClickHouse Performance.  
<https://www.youtube.com/watch?v=BHclEszF6Fk>
7. **Supercharge Your ClickHouse Data Loads - Part 1**  
ClickHouse Blog. (2023). Supercharge Your ClickHouse Data Loads - Part 1.  
<https://clickhouse.com/blog/supercharge-your-clickhouse-data-loads-part1>

